

A Combinatorial Puzzle from Arithmetic

In Praise of Recreational Mathematics

Kedar Mhaswade

This article discusses an algorithmic approach and a computer program based on it to solve an elementary combinatorial problem. Its intended audience is high school students (and laypeople) with an aptitude for abstract thinking and an interest in problem-solving. The algorithm presented here is correct but inefficient. However, it sheds light on a well-known problem in theoretical computer science.

1. Introduction

Consider the following puzzle¹:

Puzzle 1. How can you make 24 using the numbers 1, 3, 4, and 6 exactly once each, and the four arithmetic operations (+, −, ×, ÷) any number of times^a?

^aUsing a number by concatenating two or more numbers (e.g., 13) is not allowed

The correct answer is probably just an Internet search² away. Just for fun, however, one should find such an arithmetic expression without searching the Web. Try it before you read further. Don't worry, the author took quite some time to find it. If you don't find it and are restless, skip to the end.

1.1 Initial Impressions

This puzzle is surprisingly challenging, especially if one is not well-versed in solving such puzzles. We take arithmetic for granted,



Kedar Mhaswade is a computer programmer and an independent researcher. He studied mechanical engineering at the Government College of Engineering in Pune, India and computer science at IIT Chennai (1997). He has worked as a computer programmer at several companies since then. He is interested in logic and foundations of computer science, computer algorithms, functional programming, and mathematics and science education. His ORCID ID is 0009 – 0004 – 0436 – 5661.

¹The author first heard it from Sunil Singh, a mathematics educator

²Or a chat with your favorite Large Language Model or agentic AI



GENERAL ARTICLE

but the four basic operations combine with four small numbers to render it tantalizingly difficult. This is so even if you could solve it.

In our attempts, we may start with heuristics: If we can make an expression using 1, 3, and 4 that evaluates to 4, then we can make 24 by multiplying that expression by 6. But now we have a different puzzle that is just as tricky.

We may try parenthesizing parts of shorter expressions and combining their results (e.g. $((6 + 4) - 1) \times 3 = 27$). After toiling for a while we might *luckily* find the magic expression.

But we may not. What if no such expression exists? Can we tell if one exists at all? If more than one solution exists, can we find them all? Can we even tell if we have examined *all* possible expressions?

We keep *feeling* inadequate as we continue searching . . .

2. An Algorithmic Approach to Searching

As we keep searching, we may feel frustrated. After all, what we repetitively do is make expressions of four small numbers and familiar arithmetic signs and evaluate them. It gets quite mechanical and hence boring after a short while. But we tend not to give up for a perceived shame grips us. After all, we have been seeing such numbers and signs since first grade!

Even if we solve the puzzle, we wonder if we weren't lucky. If we are about to abandon³ the search and declare that no solution exists, we seem to lack the confidence that we have indeed examined every expression.

Most of all, we feel the need for assistance from a machine that accepts (or, better yet, constructs) expressions and evaluates them correctly and, well, mechanically. It's a solace that just such an accomplice exists: A computer!

³Indeed, a friend, after trying for a while, hesitantly asked the author, "Do you know if a solution exists?"



GENERAL ARTICLE

A computer is astoundingly fast at performing known calculations precisely and, at the same time, naturally incapable of possessing intelligence⁴ or originality in devising *abstractions*, the basis of computing. Alan Perlis once said [4] that, like a few other things, a computer, by definition, is *explicitly programmed* using clear instructions.

Clear instructions are called an *algorithm*. We need an algorithm to solve every problem we delegate to the computer. Perhaps surprisingly, however, the idea of an algorithm is independent of a physical computer. According to Donald Knuth [1], Euclid described the first “computer algorithm”, known as Euclid’s GCD algorithm, more than 2200 years before the advent of a digital computer!

So, how do we fit this puzzle within the confines of the strict formalism that a computer demands? If only we could construct every expression ...

2.1 Some Formal Background

To develop a perspective, we need to precisely define *expression*, *operator*, and *operand*.

An operator has a symbol and describes an operation on several operands which are just numbers in our case. Every operation *takes* some operands and *produces* a result (or results) that is, in many ways, like its operands—again, numbers in our case. Depending upon how many operands the operation takes, an operator can be classified *unary* (takes one operand), *binary* (takes two operands), etc. This property of an operator is called its *arity*. Operators can also be *commutative* or *non-commutative*. Results of commutative binary operators do not depend on the *order* of their operands: $a + b = b + a$, whereas those of non-commutative binary operators do: $a \div b \neq b \div a$

We learn in grade school that unary operator symbols are written either before the numbers (e.g. $\sqrt{8}$ to denote the square root of 8)

⁴AI techniques claim to rectify that limitation



GENERAL ARTICLE

or after the numbers (e.g. $3!$ to denote the factorial of 3), whereas the binary operator symbols are written between them ($2 + 3 = 5$). For binary operators, such an expression is called an *infix* expression.

The four basic arithmetic operations—addition, subtraction, multiplication, and division—are all binary. They take two numbers and produce a number. That helps us combine them sequentially to form expressions. Each expression can be *evaluated* to produce a value which is a number in our case.

Combining operations in expressions such as $2 + 3 + 8 + 7$ achieve brevity⁵ which aids comprehension. However, ambiguity arises when different operators appear in the same expression (e.g. $1 + 2 \times 3$) because we don't know whether to add first and then multiply (producing 9) or vice versa (producing 7). To address the ambiguity, we could have made a rule: *Do the operations in the order of their appearance.*

Instead, we use parentheses to disambiguate: $(1 + 2) \times 3$ evaluates to 9 and $1 + (2 \times 3)$ to 7. If we do not want to introduce parentheses, however, we need to define *precedence* of operators. This is where the mnemonics like “BODMAS” originate. Infix notation is a rules-heavy way to denote and evaluate expressions. But is that the only way?

Keywords

algorithms, arithmetic, recreational mathematics, combinatorics, reverse polish notation

2.2 Postfix Expressions and the Stack

As students, we should develop the habit of *meaningfully challenging the status quo*. History of mathematics and other sciences teaches us about people who dared to explore and we learn a great deal about trailblazing from them.

The Polish logician Jan Łukasiewicz⁶ proposed a different idea to denote expressions [2]:

Place the symbol after (or before) the operands.

⁵Imagine the plight of reading $2 + 3; 5 + 8; 13 + 7$ instead

⁶Pronounced *yahn woo-ka-shay-vitch*



GENERAL ARTICLE

We call such an expression a *postfix*⁷ (or *prefix*) expression. Thus, the postfix expression corresponding to the first expression above is $1\ 2\ +\ 3\times$ and that to the second is $1\ 2\ 3\ \times\ +$.

To evaluate expressions, we propose the idea of a *stack*. A stack is an abstract structure that resembles a stack of dishes we often see at restaurants. If we only allow adding and removing a dish to and from the top of the stack, then we must first remove a dish added last. This view makes stack a linear queue whose elements are added and removed only from one of its two ends—the top⁸.

Two operations on a stack are essential: 1) `push element`, which pushes the given element on its top, and 2) `pop` which simply removes the element at its top, letting us examine it. Doing the `push` operation once increments the *size* of the stack whereas doing the `pop` operation once decrements it.

If you have never encountered a stack before, you may not be impressed with the immense power its simplicity offers to *model* in many practical situations. But that is the usefulness of abstractions: It doesn't matter to the stack what it is employed to store (dishes, objects, or numbers). It provides a structure to store it and the two essential operations: `push` and `pop` defined above. Optionally, we may want to know if the stack is empty (in which case we can't pop an element) or the stack is full (which can happen in practice where resources are limited and we can't push an element).

Perhaps you are thinking of postfix expressions as an unnecessary digression. You may be wondering why a postfix expression like $1\ 2\ 3\ \times\ +$ is more convenient than the corresponding infix expression: $1 + (2 \times 3)$.

And one might even successfully argue that infix expressions are easier to read. But isn't readability a question of getting used to? Had we been reading postfix expressions since grade school, we might have become more comfortable with them. Is there another benefit?

⁷It's also called the *Reverse Polish notation*

⁸A stack is therefore called a LIFO—Last In First Out—queue



It turns out that postfix expressions need no operator precedence rules, and, as a result, they do not need parenthesizing expressions to enforce precedence. A given postfix expression can be evaluated *mechanically* (with the help of the stack)! This bonus becomes apparent when we study the *algorithmic* nature of evaluating postfix expressions.

We start with a postfix expression. For the sake of our algorithm, we read it a *token* at a time. A token is either an operand or an operator. Our operators are all binary. However, the scheme can be easily extended to other operators (e.g., unary). We use the stack only to store operands and results of operations. Thus, ours is a stack of numbers. Here is our Algorithm (1).

Algorithm 1 Evaluate any postfix expression with binary operators

```

1:  $S \leftarrow \text{STACK.NEW}$                                  $\triangleright$  Make a new stack
2:  $T \leftarrow [t_1, t_2, \dots, t_n]$ 
3: while T has tokens do                                 $\triangleright$  Scan the expression from left
4:    $t \leftarrow T.\text{NEXT}$                                   $\triangleright$  Get the next token
5:   if t is an operand then
6:     S.PUSH(t)
7:   else if t is an operator then     $\triangleright$  Write defensive program
8:     od2  $\leftarrow S.\text{POP}()$            $\triangleright$  Pop two operands
9:     od1  $\leftarrow S.\text{POP}()$            $\triangleright$  Order matters!
10:    Return Error: not enough operands on the stack
11:    r  $\leftarrow \text{OP}(od1, od2)$          $\triangleright$  Do it: e.g.  $2 + 3$  and  $r \leftarrow 5$ 
12:    S.PUSH(r)
13:   else
14:     return Error                   $\triangleright$  Impossible? An invalid token
15:   if S.EMPTY( ) then
16:     return Error                   $\triangleright$  S must have the result
17:   else
18:     return S.POP()  $\triangleright$  A valid expression evaluates to this value

```



2.3 A Helpful Insight?

Infix and postfix (or prefix) are, after all, just notation. And it so happens that one can deterministically translate a valid infix expression into postfix and vice versa. However, removing the need for parentheses and providing a mechanical way of evaluating are strengths of the postfix expression which the author used to address our puzzle (1). Let's list the helpful facts that we know so far. They may sound disparate at the moment, but they might align themselves in a coherent whole later:

We want to employ a (preferably systematic) method to *assuredly* examine all expressions made of exactly four numbers (1, 3, 4, 6) and required (binary) arithmetic operators (+, −, ×, ÷). The emphasis is on not missing any valid expression.

Can postfix notation help?

Since there are four operands and each *binary* operator takes two operands and produces a result (which may then pair up with the next operand), we must have exactly three binary operators in an expression with seven tokens.

We have four operands that can appear in any order in a postfix expression. Since we must have exactly three binary operators to operate on them, there are seven places each of which is occupied either by an operator or an operand.

Here is a possible arrangement:

4	1	−	3	+	6	÷
---	---	---	---	---	---	---

Fig 1: A Valid Arrangement

Luckily, this is a valid postfix expression that evaluates to 1. We can trace the run of our algorithm (1) on this expression to confirm the result.



Here is another possible arrangement:

-	1	3	4	+	6	÷
---	---	---	---	---	---	---

Fig 2: A Valid Arrangement, but Invalid Postfix Expression

We are unlucky in this case because this arrangement is not a valid postfix expression; it does not evaluate to any value. Expectedly, our algorithm (1) returns an error in this case. But a computer doesn't care and perhaps neither should we, because we can ignore this arrangement as a “non-expression” and examine the next one. It would be nice to know *a priori* that such an expression is invalid (without actually evaluating it with our algorithm (1)), but that is not necessary. The human mind may *immediately* discard all the expressions like (2) that begin with an operator (because there is nothing to operate it on).

2.4 The Algorithm

Perhaps we now feel closer to describing a foolproof⁹, if inefficient, method of solving our puzzle to the computer. Our search has got to be exhaustive; **we must not miss any (valid) arrangement of four operands and three binary operators** but it is acceptable if some arrangement leads to an invalid expression. We evaluate each expression. If an expression evaluates to our target number (24), we have succeeded. If there is no such expression, the puzzle is unsolvable. An algorithm seems to emerge. We just need to describe the process precisely ...

We are given a set of operands, but where they appear in the seven locations (their *order*) matters. This is an archetypal situation leading us to a fascinating area of mathematics called Combinatorics. Specifically, we can arrange k objects in n locations¹⁰ in $\frac{n!}{(n-k)!}$ ways. In our case, the number of such *ordered arrangements* of four operands in seven locations is $\frac{7!}{3!} = 7 \times 6 \times 5 \times 4 =$

⁹Being foolproof is required because a computer is going to carry it out

¹⁰Each such arrangement is called a *Permutation*



GENERAL ARTICLE

840. Here are some of the arrangements:

1	3	4	6				
1	3	6	4				
...
1	4		3		6		
...
		1	3		6	4	
...
			6	4	3	1	

Fig 3: All Permutations of Four in Seven

These are *all* the possible ways in which we can arrange members of the set of operands {1, 3, 4, 6} in seven places, and there are exactly 840 such arrangements. Some arrangements will invariably lead to invalid postfix expressions (e.g., the permutations in red in Figure (3)) after we place the operator symbols in the remaining places.

How do we place operator symbols? Well, we have three places and we want to place three operators in them. However, this time repetition is allowed because we can choose the same operator more than once. Such arrangements where members of an n -set are placed in k places (where choosing a set member more than once is allowed) are called k -tuples [5] of the n -set. The number of k -tuples is n^k . In our case $k = 3, n = 4, n^k = 64$. We have 64 3-tuples of 4 operators to consider.

Let's demonstrate the placement of all operator 3-tuples with one arrangement of operands, [1,4,_,3,_,6,_], (shown in green) from Figure (3).

1	4	+	3	+	6	+	
1	4	+	3	+	6	-	

 $= 14$

 $= 2$

GENERAL ARTICLE

1 4 + 3 + 6 ×	= 48
1 4 + 3 + 6 ÷	= 1. <u>333</u>
1 4 - 3 + 6 +	= 6
1 4 - 3 + 6 -	= -6
...
1 4 ÷ 3 ÷ 6 ×	= 0.5
1 4 ÷ 3 ÷ 6 ÷	= 0.13 <u>888</u>
...
+ - 1 4 × 3 6	= Error

Fig 4: All Postfix Expressions for the Operand Permutation $[1, 4, 3, 6]$

Consider a valid arrangement from Figure (4):

1 4 + 3 + 6 ×

Here is how the Algorithm (1) behaves in that case (scanning proceeds from left to right; stack, initially empty, grows up):

Token	Lines from Listing (1)	Stack	Tokens Yet to Scan
1	5–6	1	$[4 + 3 + 6 \times]$
4	5–6	4 1	$[+ 3 + 6 \times]$
+	7–11	5	$[3 + 6 \times]$
3	5–6	3 5	$[+ 6 \times]$
+	7–11	8	$[6 \times]$
6	5–6	6 8	$[\times]$



GENERAL ARTICLE

×	7–11	48	[]
	16–17		

The expected result (48) is returned.

An upper bound on all 7-token postfix expressions thus formed is: $64 \times 840 = 53760$. We examine every expression, and if it is valid, check the result it produces. Some *optimizations* can reduce the number of expressions we need to evaluate. We assume the availability¹¹ of procedures `PERMS`, `TUPLES`, `EXPR`. The `EVALPOSTFIX` procedure calls Algorithm (1).

The listing for the algorithm that performs the exhaustive search is at (2). It is generic and therefore useful in practice. Of course,

Algorithm 2 An Algorithm to Solve Puzzle (1)

```

g ← 24                                ▷ Target result: 24
2: f = false                            ▷ Did we find an expression?
P ← PERMS()                           ▷ P ← [P1, P2, ..., P840]
4: T ← TUPLES()                          ▷ T ← [t1, t2, ..., t64]
   for all p ∈ P do
6:     for all t ∈ T do
      e ← EXPR(p, t)    ▷ e ← from given permutation, tuple
8:     r ← EVALPOSTFIX(e)          ▷ r ← use algorithm (1)
      if r = g then output Found g: e
10:    f = true
      if f = false then
12:        output There's no way to make g!

```

many optimizations are possible. It is also extensible to unary operators such as the factorial. However, unary operators cause ambiguity because, unlike binary operators, they can be applied any number of times. For instance, if we use the unary $\sqrt{}$ (square root) operator, the expression: [16, $\sqrt{}$, $\sqrt{}$] makes 2 from 16.

¹¹We implement them at (2.5)



2.5 An Implementation

A Java implementation of the algorithm (2) is available at [3]. It correctly produces the expressions that solve this and similar problems.

3. Summary and Conclusion

The author describes a surprising application of a technique he learned in graduate school to solve a combinatorial puzzle he encountered on social media years later. Of course, there may be other techniques to solve the puzzle.

The puzzle is simple enough for any curious layperson familiar with basic arithmetic to attempt and enjoy, but it begets some serious computational questions.

1. Every possible expression must be examined. There is no “shortcut” to find the expression that evaluates to a target number. Must we examine every expression? Can we do *better* than the mundane “make and evaluate every expression” method?
2. If someone claims to find an expression, however, we can efficiently verify if the claim is valid.
3. For four operands, we have to analyze more than 50000 expressions in the worst case. How many expressions might we have to analyze if the number of operands is 10? Or 20?
4. If no expression evaluates to a given target, is there a trick that lets us *quickly* conclude that without having to analyze every expression? In other words, do we have a more efficient algorithm to solve this puzzle in the negative case?

Here is another similar puzzle:

Puzzle 2. *How can you make 74 using the numbers 1, 3, 4, and 6 exactly once each, and the four arithmetic operations (+, −, ×, ÷) any number of times^a?*



GENERAL ARTICLE

(Hint: The expression $((4 \times 6) + 1) \times 3$ evaluates to 75.)

¹²Using a number formed by concatenating two or more numbers (e.g. 13) is not allowed

The puzzle represents one class of problems that pose challenges even to the fastest computers because an “efficient” algorithm may not exist¹². Increasing the number of operands by an additive constant increases the number of cases to examine by a multiplicative constant. Who knows, if you can find an efficient algorithm for such problems or *prove* that none exists, you may win a million dollars¹³!

4. Answer to the Original Puzzle

Time for a handstand, maybe?

$$1. 6 \div (1 - (3 \div 4))$$

Acknowledgments

The author wishes to thank Deepa Joshi and two anonymous reviewers, who read the entire manuscript and verified the correctness of given algorithms.

Suggested Reading

- [1] Donald Knuth. **The Art of Computer Programming, Fundamental Algorithms.** Vol. 1. Addison-Wesley Professional, 1997
- [2] Hamblin C. L. “Translation to and from Polish Notation”. *The Computer Journal* 5.3 (1962), pp. 210–213. doi: 10.1093/comjnl/5.3.210
- [3] A Java Implementation for the Problem 24. 2025. [Fazlur](#) (visited on 02/13/2025)
- [4] Alan Perlis. **Foreword to Structure and Interpretation of Computer Programs.** 1984. (visited on 02/09/2025)
- [5] Permutation. Feb. 12, 2025. [k-tuples](#)

¹²Try asking puzzle (2) to an LLM like ChatGPT; you may not conclude that it is *intelligent*

¹³See the [P versus NP Problem](#)

